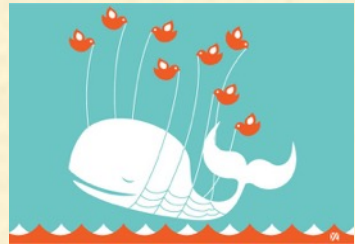

PELIKAN “CACHE À LA CARTE”

A framework for building *production-ready* cache in datacenters.

Yao Yue, Twitter Inc
@thinkingfish

ABOUT ME

- at Twitter since 2010 
 - working on cache this whole time
 - worn every hat
-

ABOUT THE TALK

- mostly *not* about cache
 - based on twitter's use of cache
 - a quest for high-quality infrastructure
-

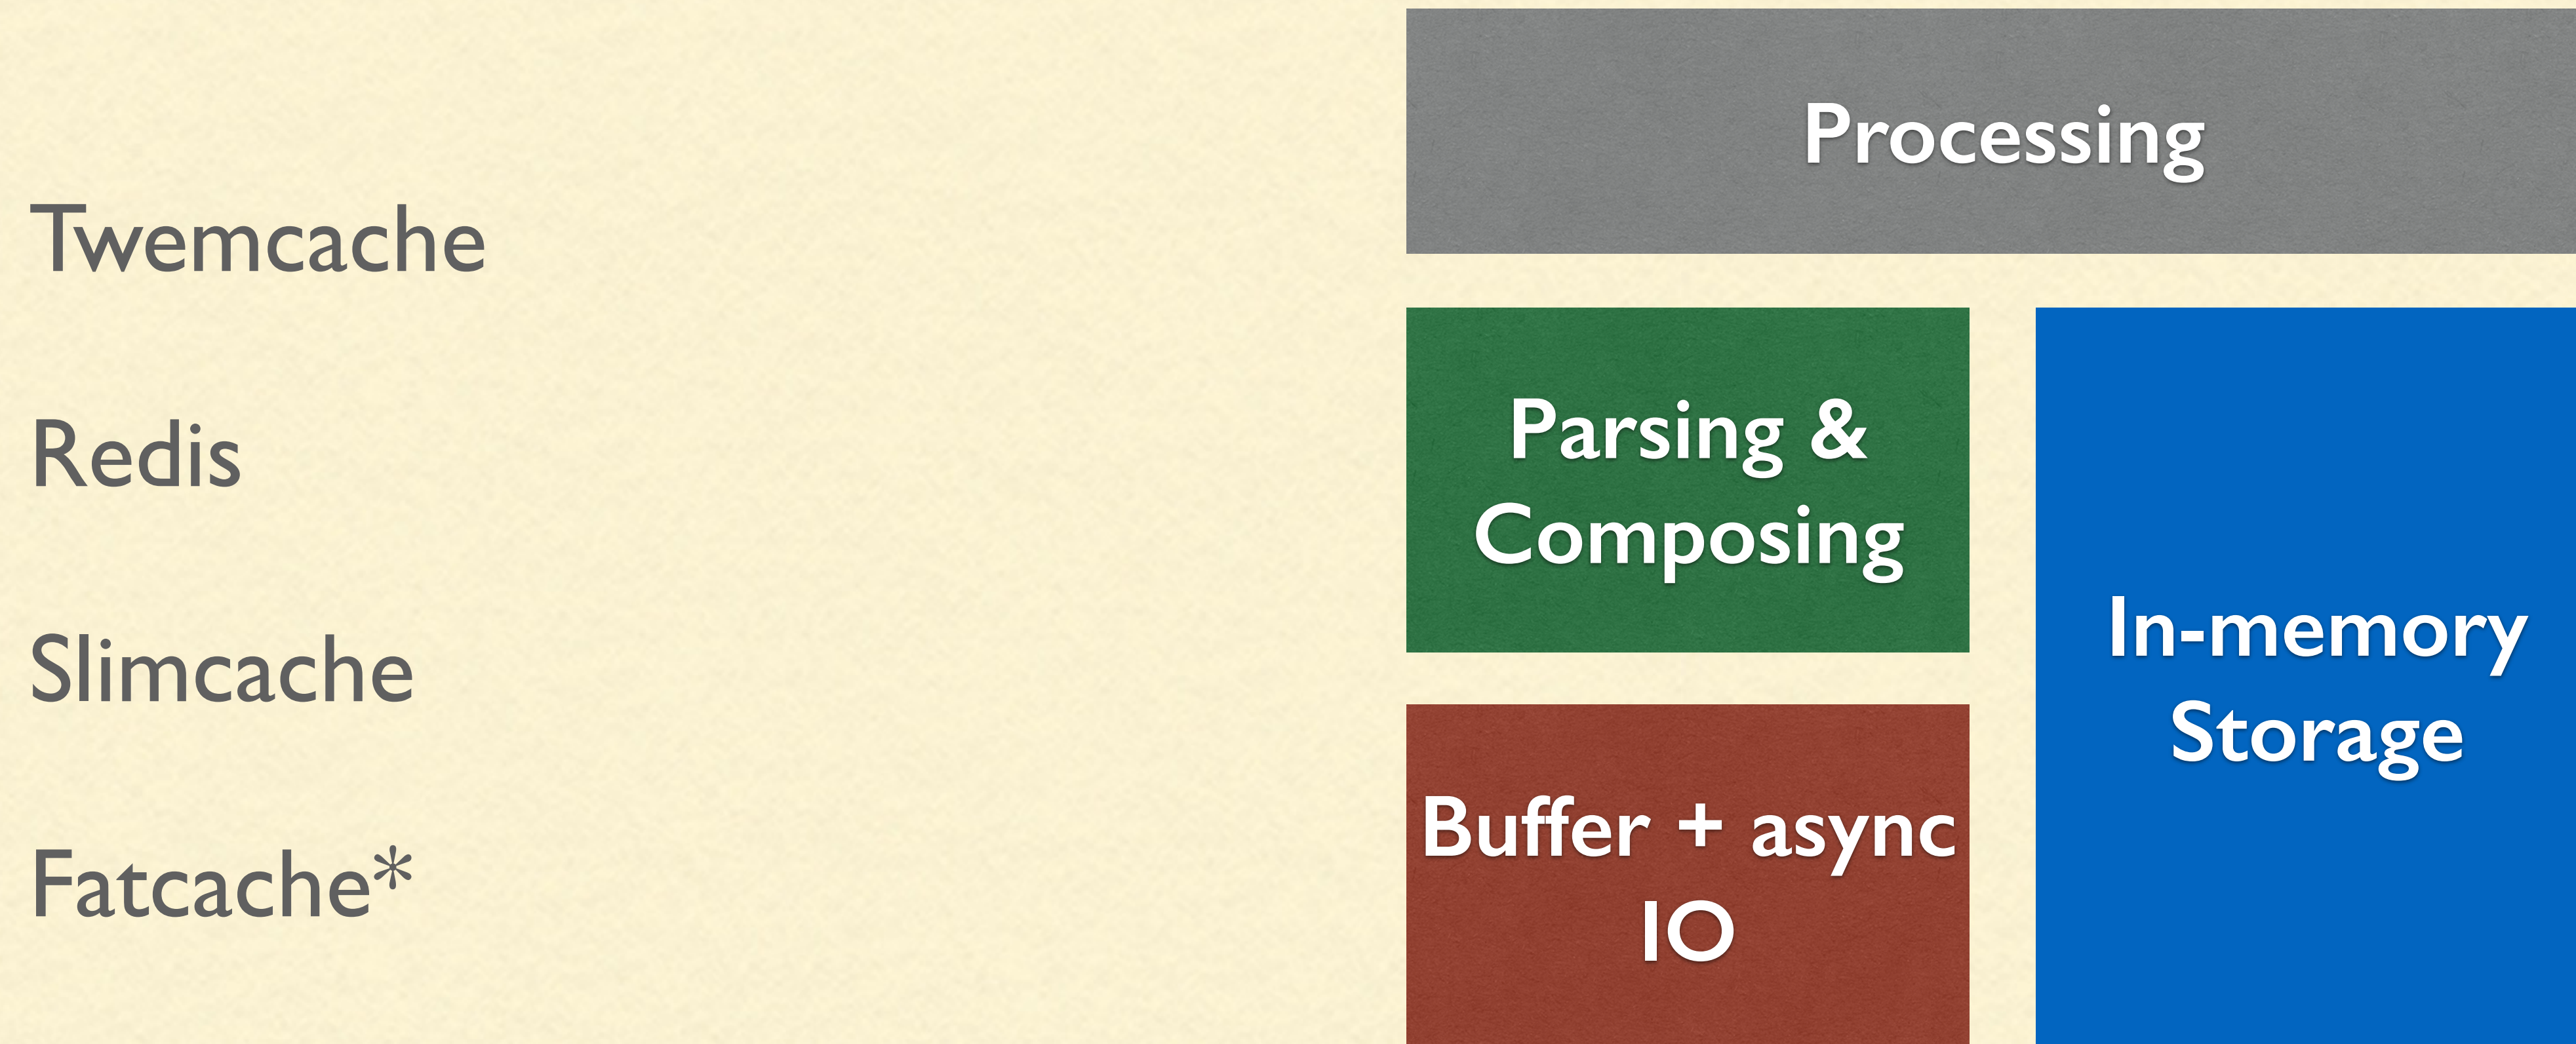
Context



DEPLOYMENT & SCALE

- many clusters, in containers, automated deploy
 - qps: from thousands to tens of millions
 - data models, query size and access pattern- all over the map
-

A “SOLVED” PROBLEM



COMMON CHALLENGES

- many ways to fall below SLA
 - hotkeys and DDoS
 - hard to debug
 - capacity planning surprises
-

THE CACHE WE WANT

- covers all our use cases
 - easy and fun to work on
 - is production-ready
-

A DIRTY LITTLE SECRET

we have little idea about what production-readiness demands

- rely on “battle-tested” solution
 - “f*ck it”
-

PRODUCTION-READY CACHE

PREDICTABLE

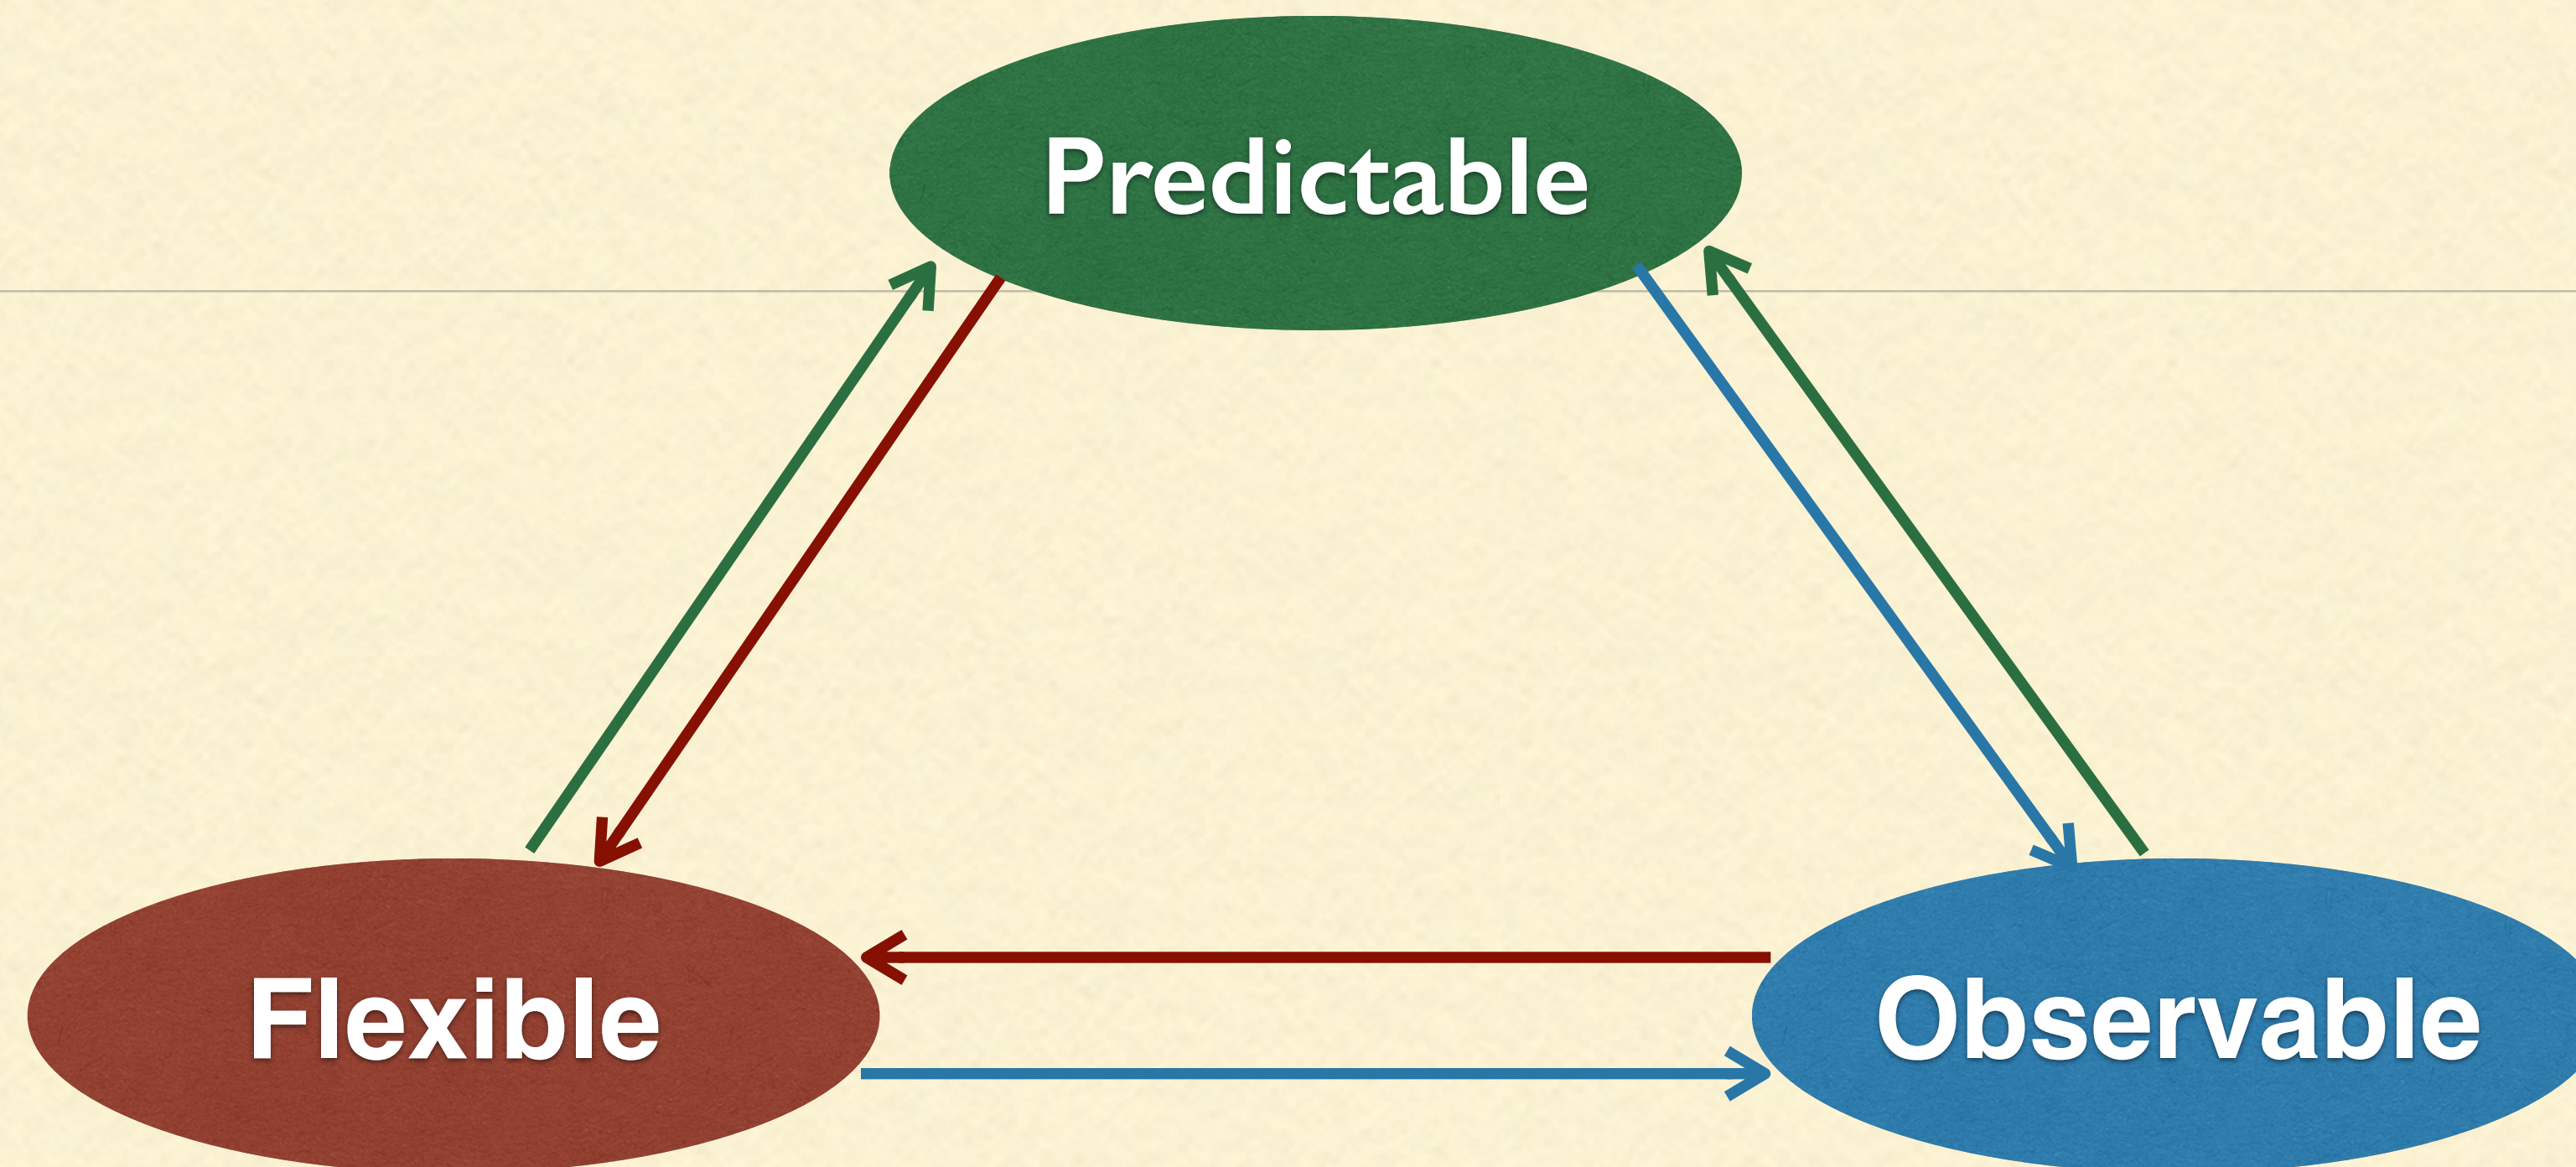
- tail-latency & performance
 - failure behavior & degradation
 - resource footprint
-

OBSERVABLE

- ready to be monitored
 - debuggable
 - reveals internal flow
 - analytics-friendly
-

FLEXIBLE

- configurable
 - composable
 - quick to develop features
-

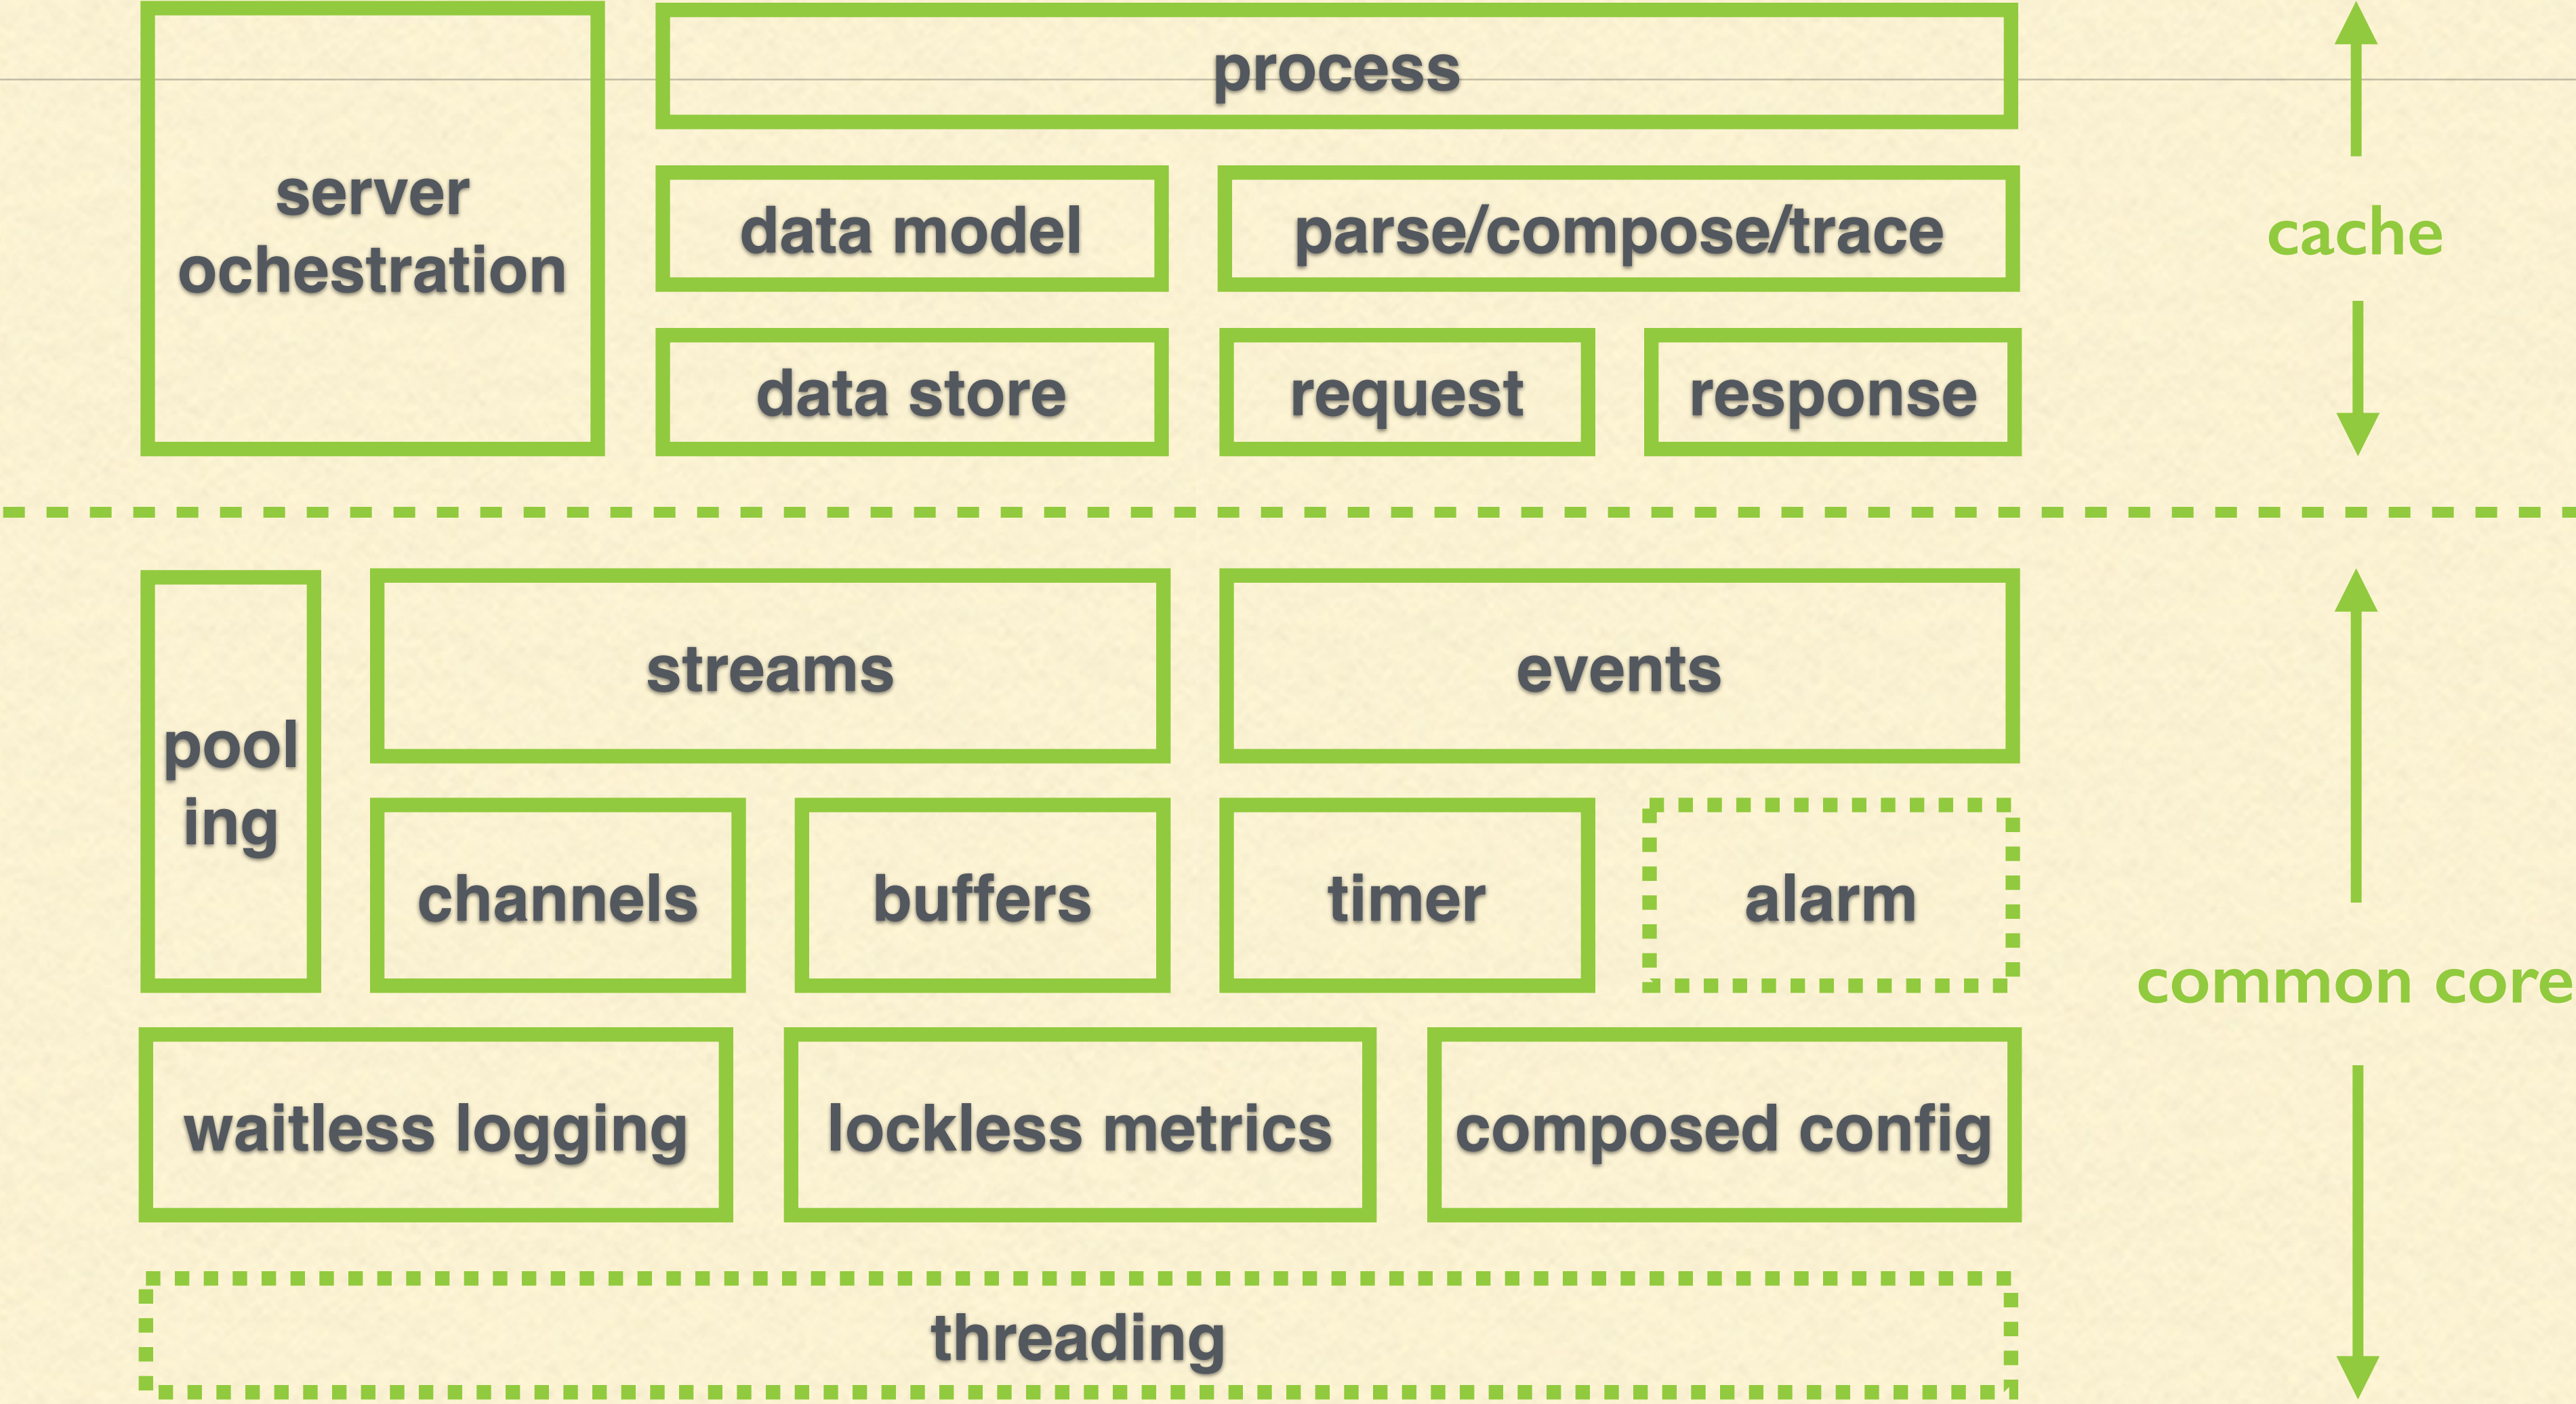


—————→
“should also be”

“SHOW US WHAT YOU GOT”



ARCHITECTURE OVERVIEW



MODULARIZE...

- minimize surface area- no leaky abstraction
 - right amount of generality
-

...AND FOR EACH MODULE:

- “Are the performance/failure scenarios known?”
 - “How do I manage the resources?”
 - “How do I get visibility? What do I track?”
 - “How can I configure this? What’s good defaults?”
-

JUDICIOUS CODE REUSE

- clean slate design
 - use tested logic
 - alter to fit
 - ~~future-proof~~ future-compatible
-

CORE DECISIONS

- common core / cache split

everything you need for a production-ready ping server + the rest

- control / data plane split

background thread performing non-critical tasks

DEVELOPMENT

- developers: me, Kevin Yang, Sagar Vemuri
 - mostly since summer 2014
 - clean-slate design v.s. ~50% existing code
 - binaries: pelikan_twemcache, pelikan_slimcache, pelikan_redis*
 - pilot production deploy done; load test**, universal canary to come
-

PRODUCTION-READINESS: SOME HIGHLIGHTS

LOG, STATS, CONFIG

ubiquitous, *paradigms*

make them cheap, configurable

make them composable

waitless logging

lockless stats

modular config

DECLARE/INITIALIZE METRICS

```
/*          name          type          description */
#define BUF_METRIC(ACTION) \
    ACTION( buf_curr,      METRIC_GUAGE,  "# buf allocated"      )\
    ACTION( buf_active,    METRIC_GUAGE,  "# buf in use/borrowed" )\
    ACTION( buf_create,    METRIC_COUNTER, "# buf creates"        )\
    ACTION( buf_create_ex, METRIC_COUNTER, "# buf create exceptions")\
    ACTION( buf_destroy,   METRIC_COUNTER, "# buf destroys"       )\
    ACTION( buf_borrow,    METRIC_COUNTER, "# buf borrows"        )\
    ACTION( buf_borrow_ex, METRIC_COUNTER, "# buf borrow exceptions")\
    ACTION( buf_return,    METRIC_COUNTER, "# buf returns"        )\
    ACTION( buf_memory,    METRIC_GAUGE,  "memory allocated to buf")

typedef struct {
    BUF_METRIC(METRIC_DECLARE)
} buf_metrics_st;

#define BUF_METRIC_INIT(_metrics) do { \
    *(_metrics) = (buf_metrics_st) { BUF_METRIC(METRIC_INIT) }; \
} while(0)
```

UPDATE METRICS

```
INCR(buf_metrics, buf_create);  
INCR(buf_metrics, buf_curr);  
INCR_N(buf_metrics, buf_memory, buf_init_size);
```


INCLUDE METRICS

```
struct glob_stats {  
    procinfo_metrics_st    procinfo_metrics;  
    event_metrics_st       event_metrics;  
    server_metrics_st      server_metrics;  
    worker_metrics_st      worker_metrics;  
    buf_metrics_st         buf_metrics;  
    tcp_metrics_st         tcp_metrics;  
    cuckoo_metrics_st      cuckoo_metrics;  
    request_metrics_st     request_metrics;  
    response_metrics_st    response_metrics;  
    parse_req_metrics_st   parse_req_metrics;  
    compose_rsp_metrics_st compose_rsp_metrics;  
    process_metrics_st     process_metrics;  
    log_metrics_st         log_metrics;  
};
```

```
struct glob_stats glob_stats;
```

```
buf_setup((uint32_t)setting.buf_init_size.val.vuint, &glob_stats.buf_metrics);
```

BUFFER, CHANNEL, STREAM

buffer connects sync/async processing

interface hides multiple implementation

all resources can pooled, capped

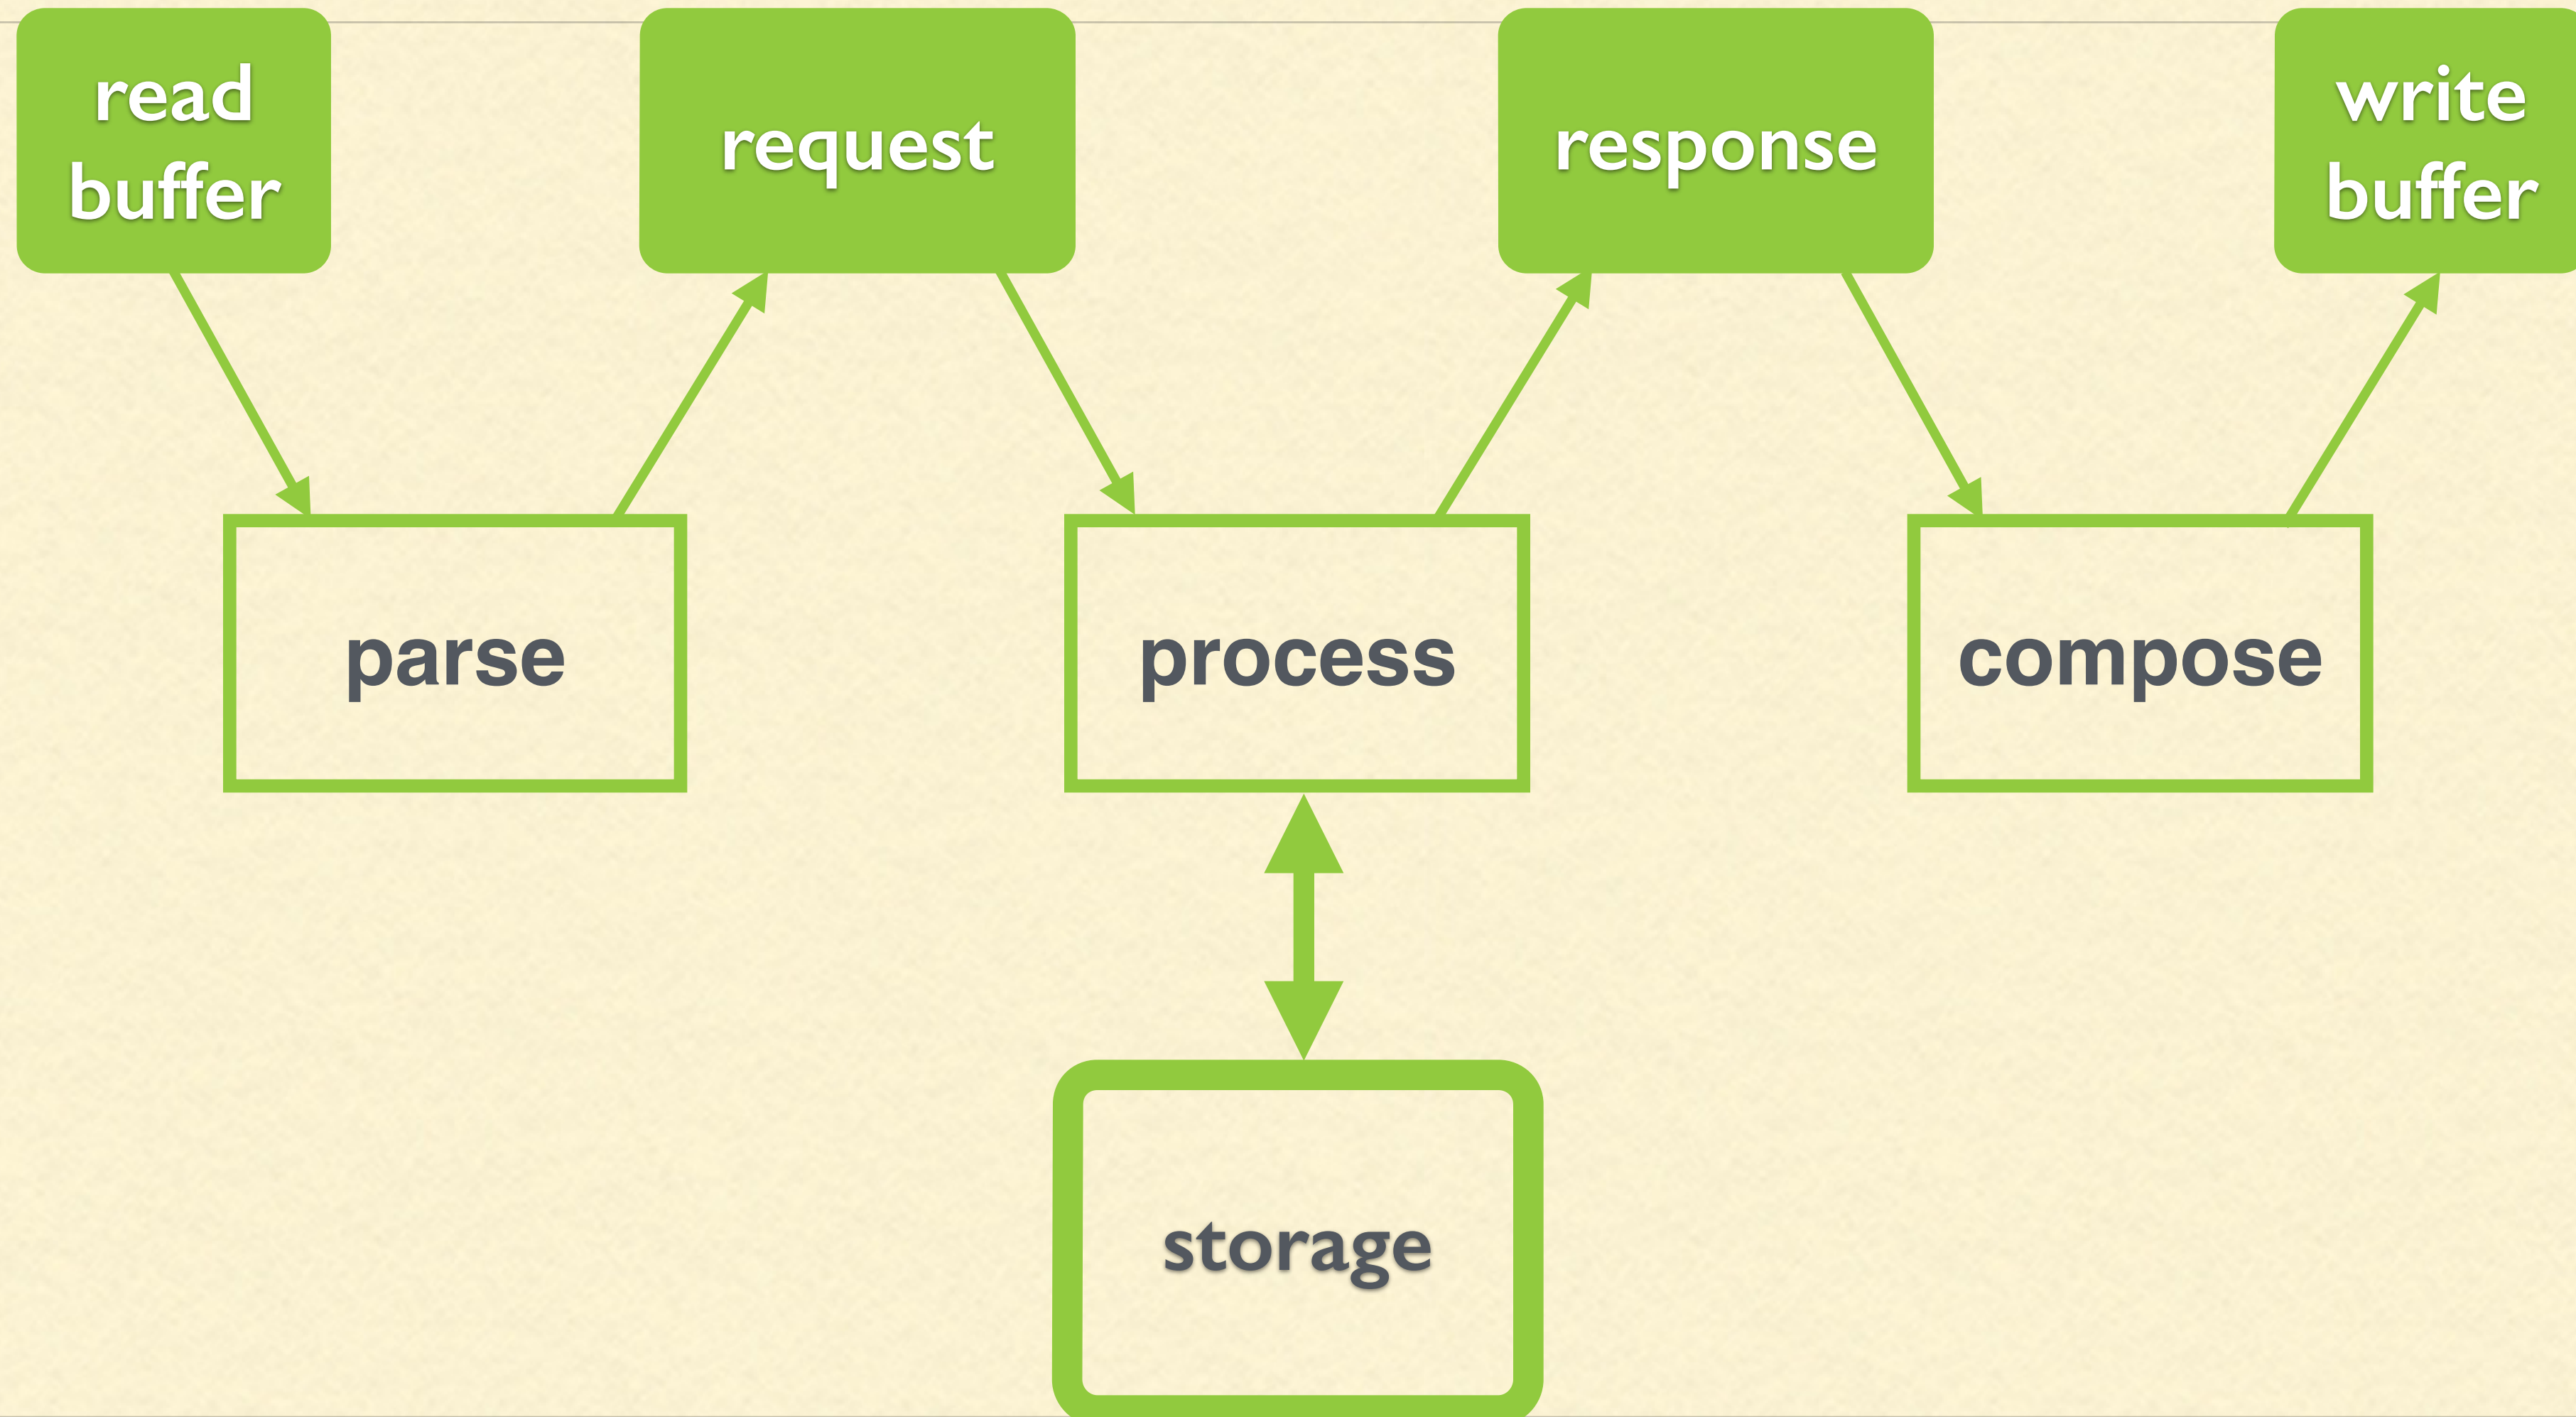
streams

buffers

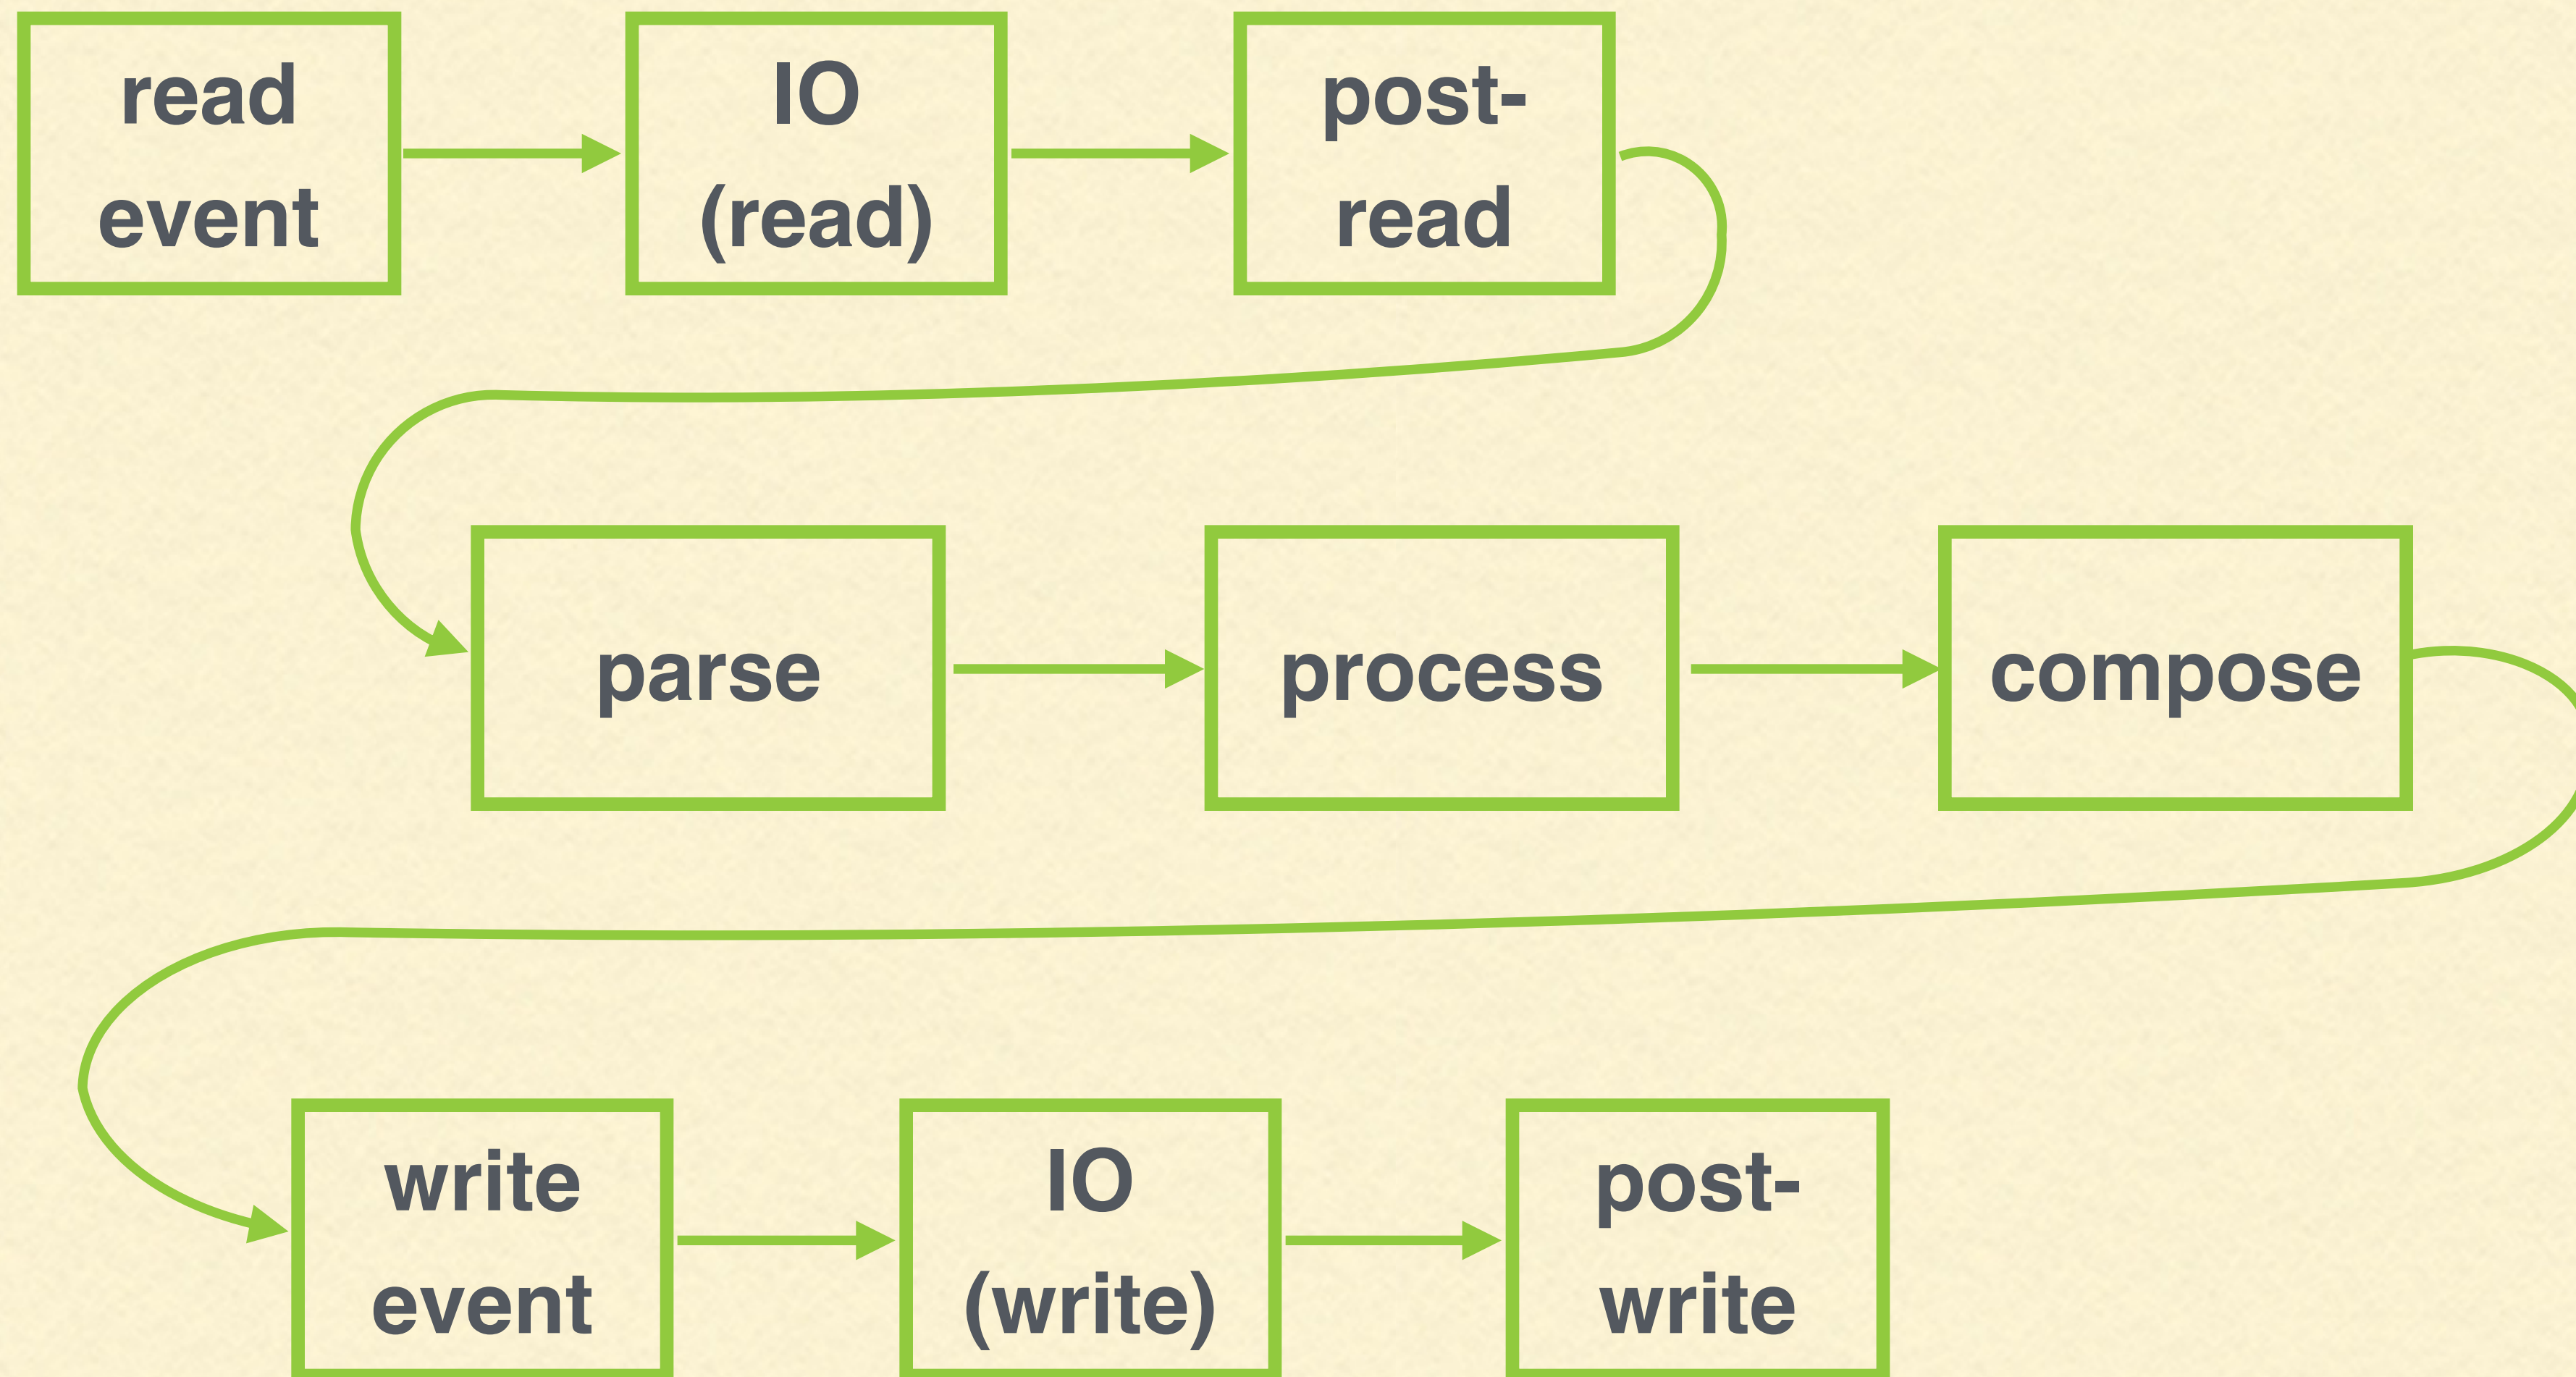
channels

**(de-)serialization
processing & other logic**

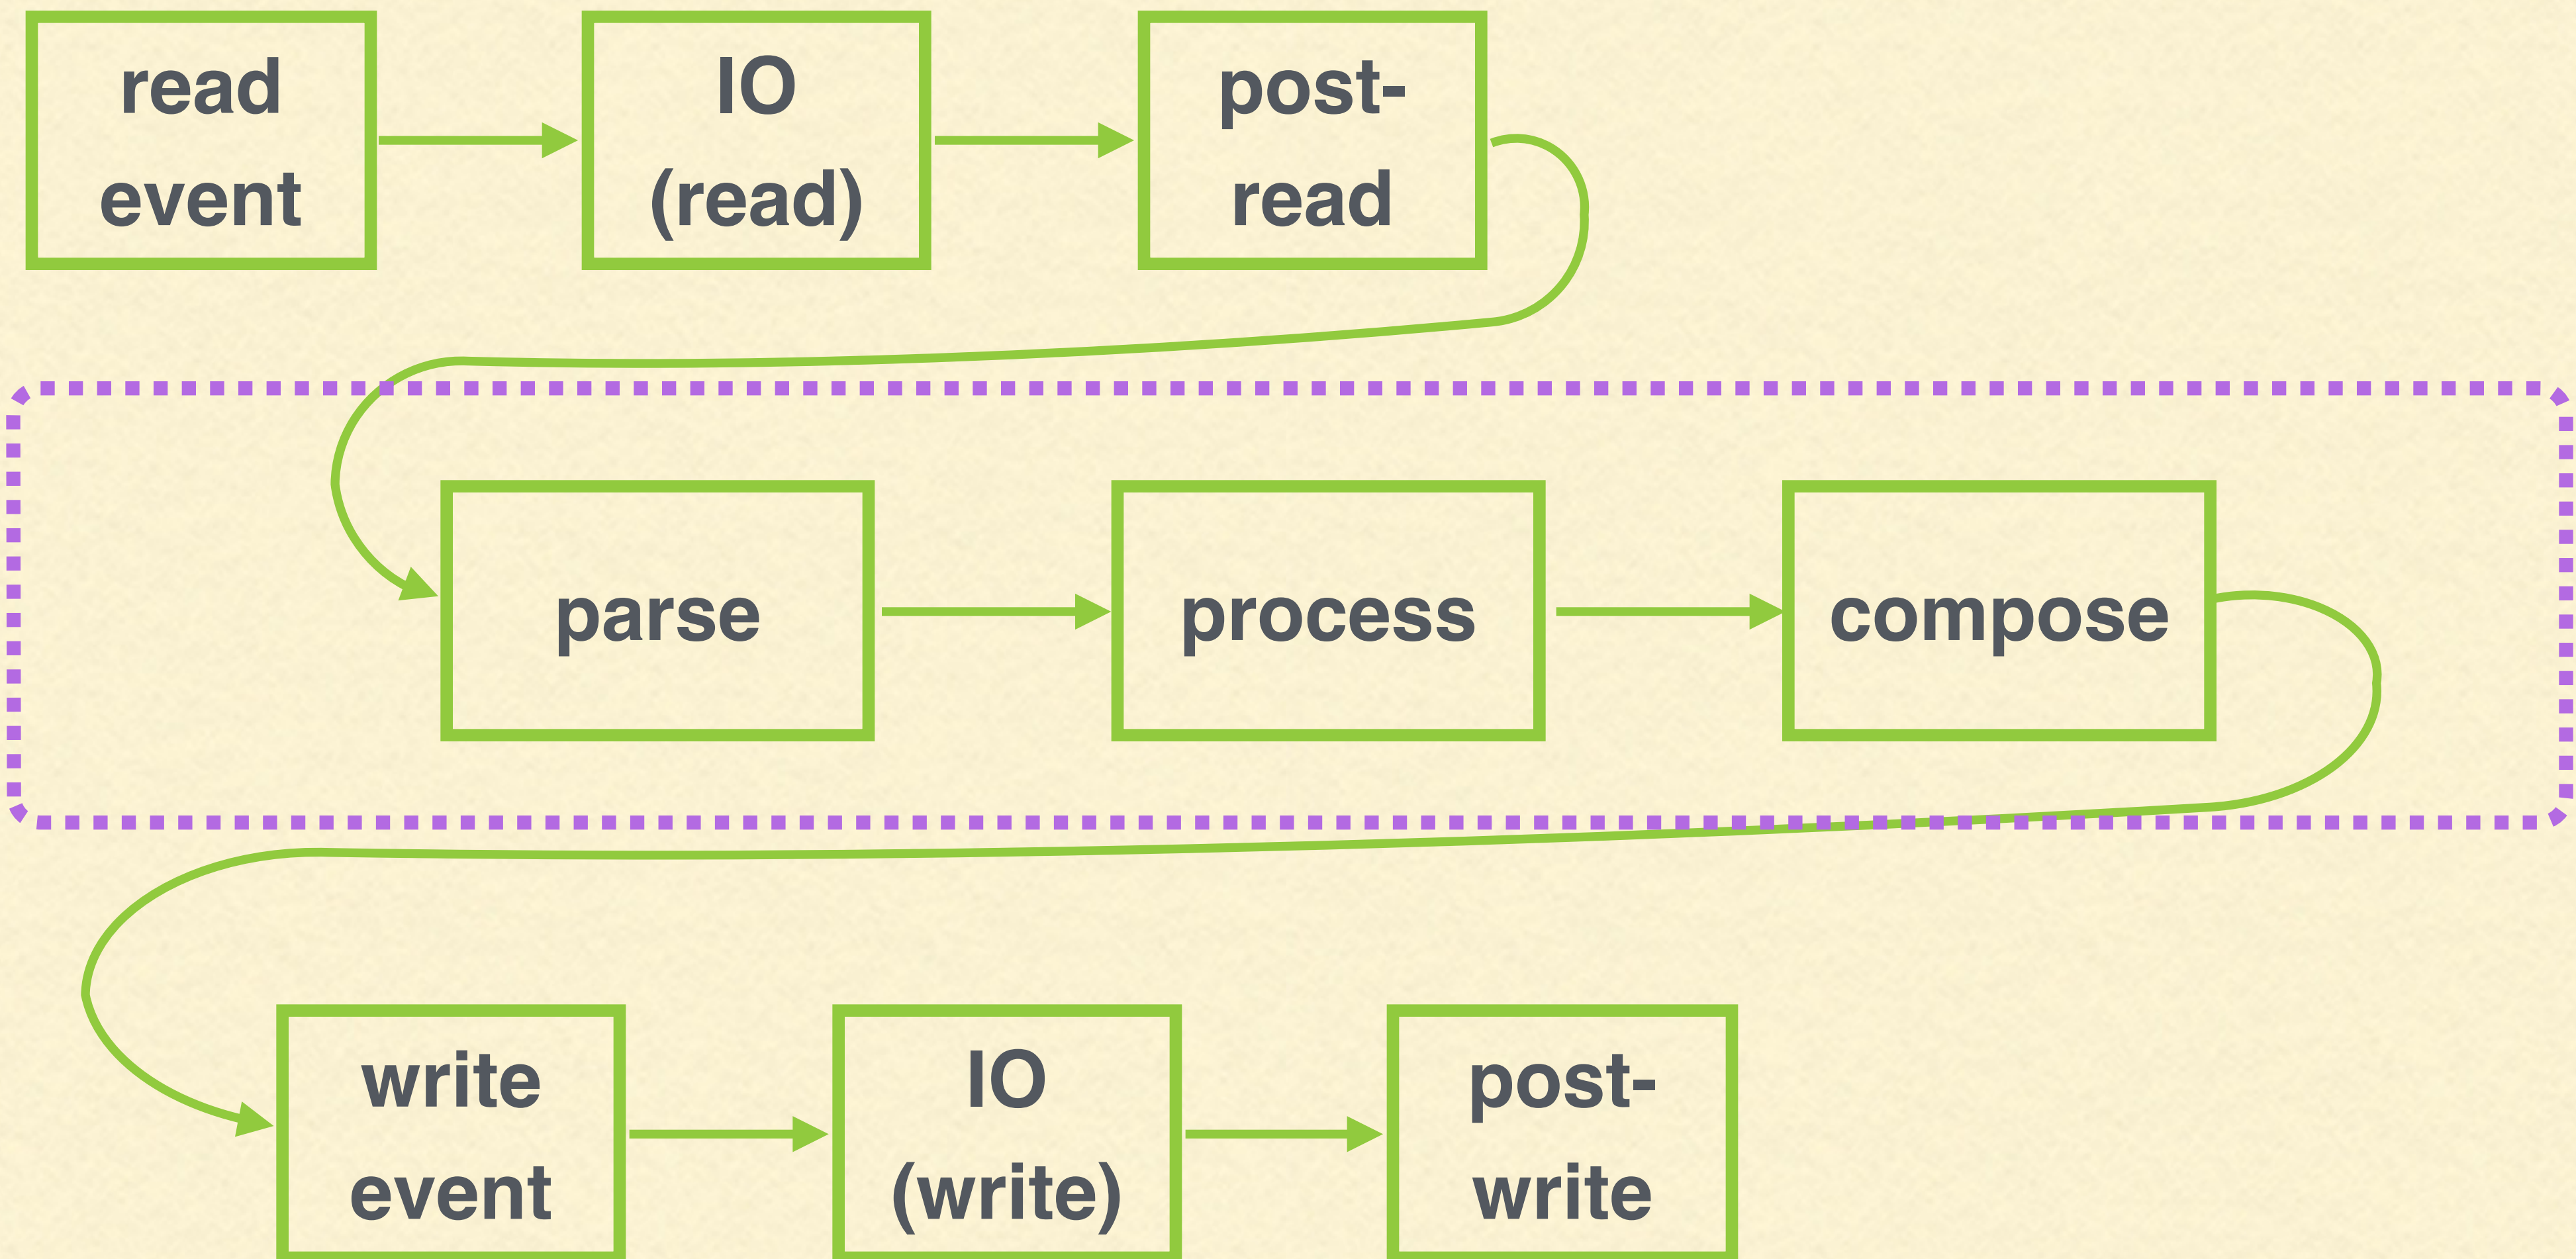
STRUCTURAL SYMMETRY



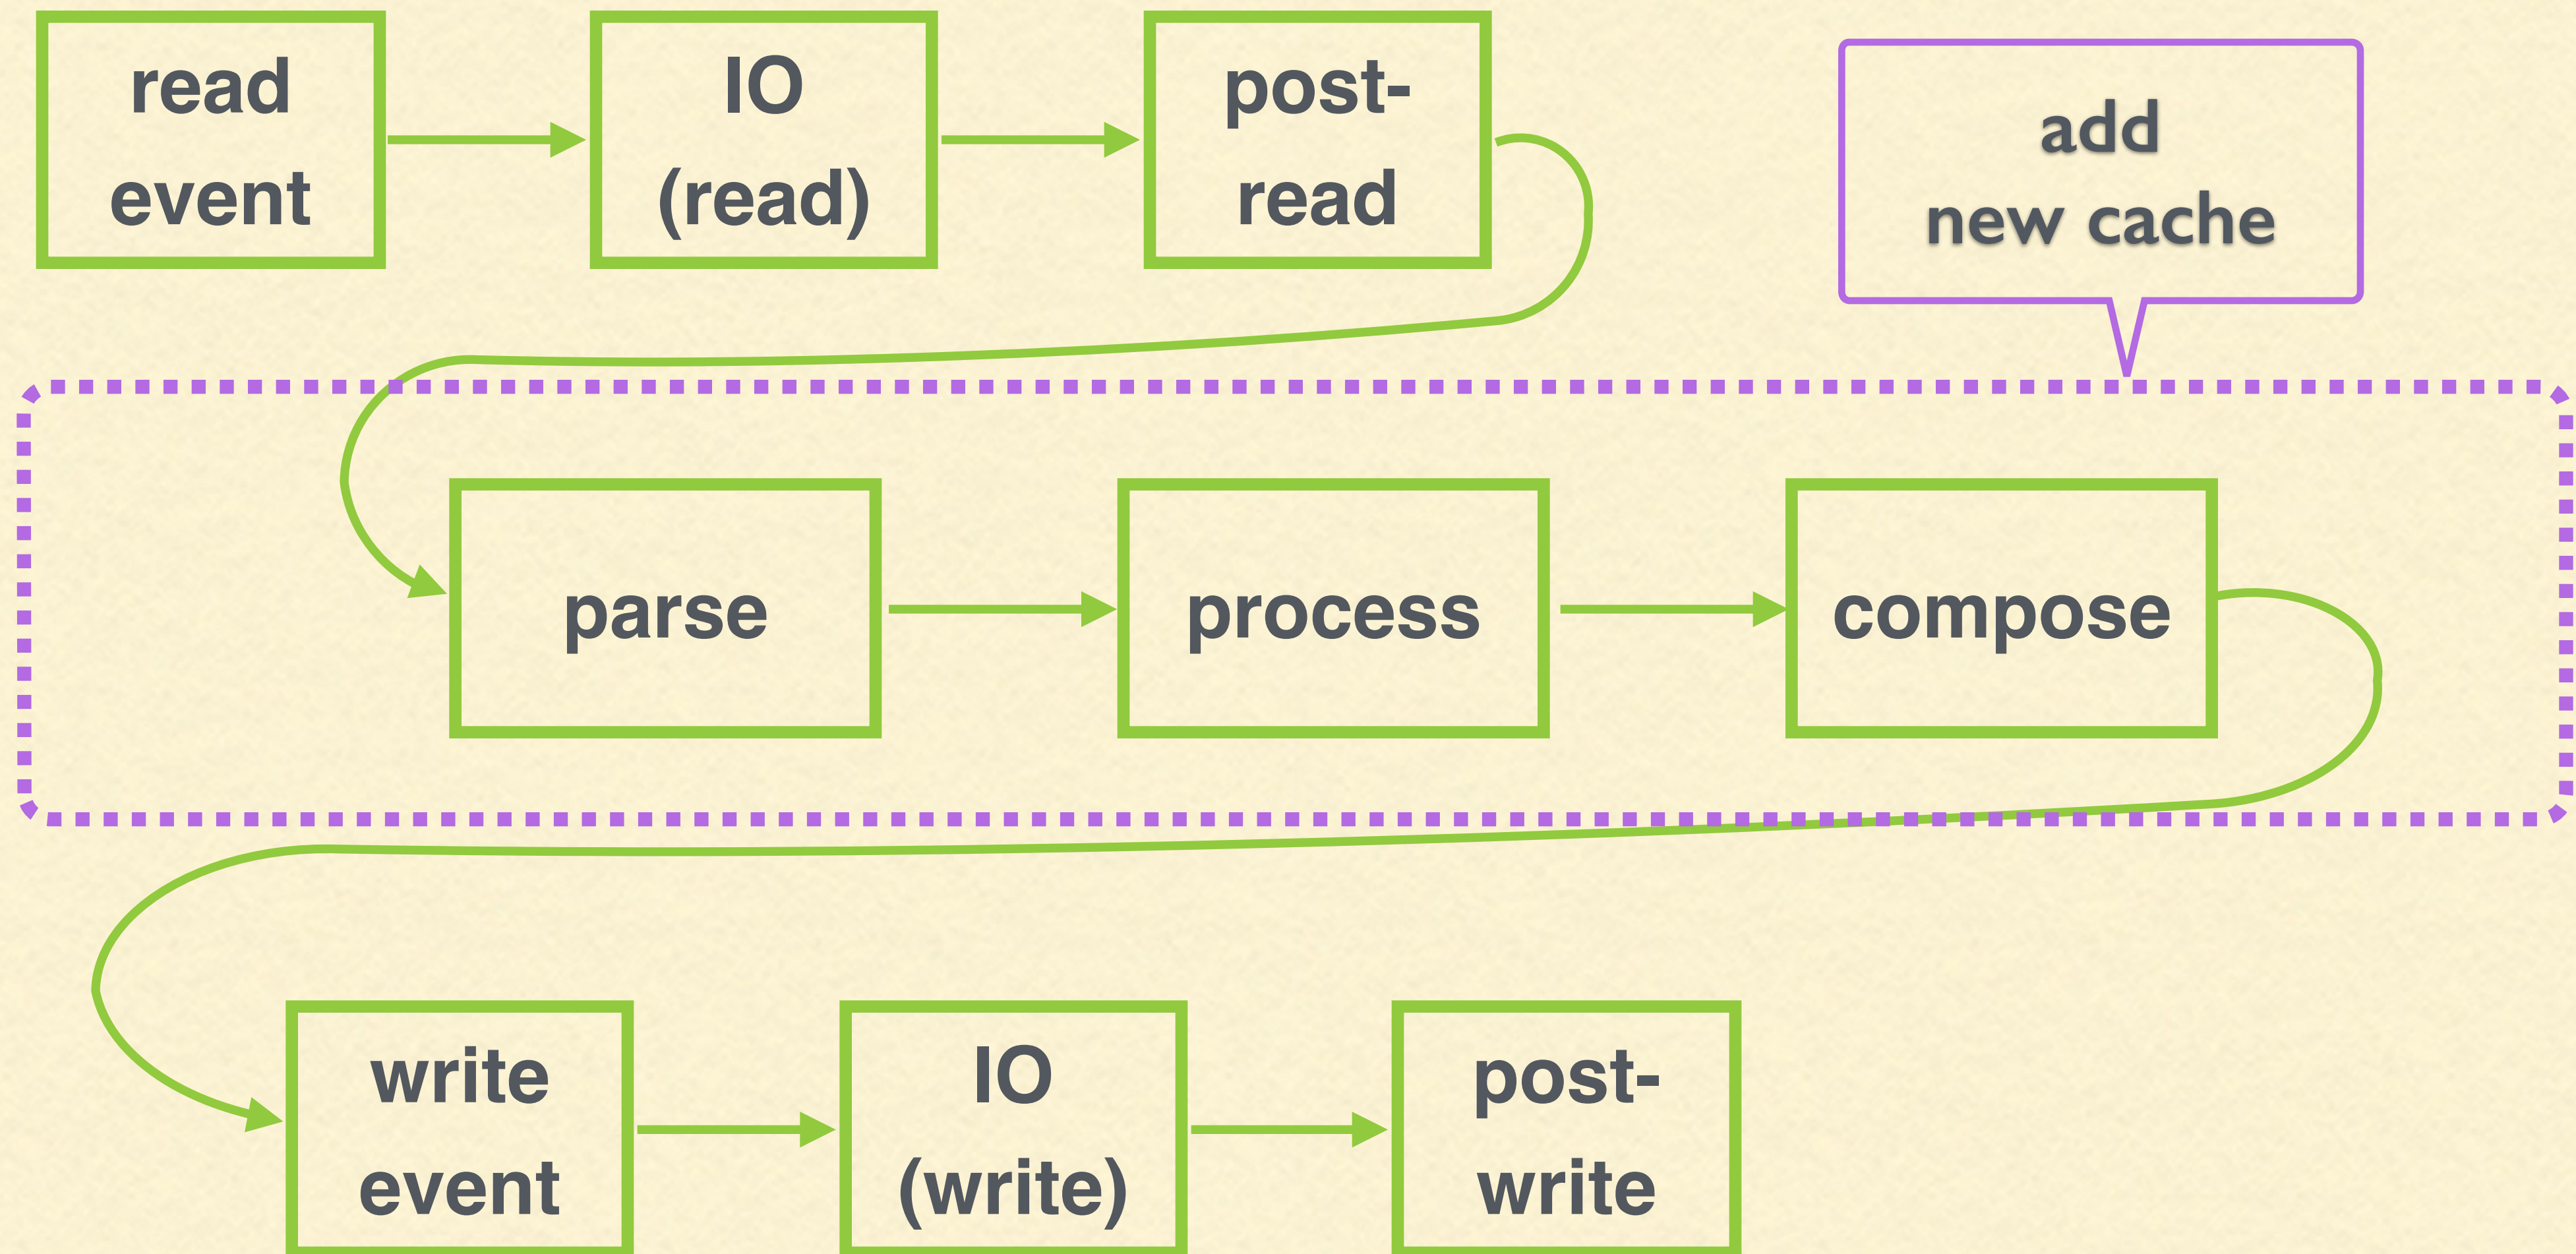
STANDARD FLOW



STANDARD FLOW



STANDARD FLOW



ADD A NEW CACHE

```
parse_rstatus_t parse_req(struct request *req, struct buf *buf);  
parse_rstatus_t parse_rsp(struct response *rsp, struct buf *buf);
```

```
int compose_req(struct buf **buf, struct request *req);  
int compose_rsp(struct buf **buf, struct response *rsp);
```

```
void process_request(struct response *rsp, struct request *req);
```

CASE STUDY: HOW MUCH CODE

- twemcache: 14k LOC

uses libevent, not counted

- pelikan_twemcache: 16K LOC


common core: 9K LOC

twemcache: 7K LOC

CASE STUDY: ADD SLIMCACHE

- goal: 6B metadata instead of ~60B per key for small objects
 - total: ~1600 additional LOC
 - cuckoo hashing for storage: 752 LOC
 - process module: 506 LOC
 - other code for a new executable: 323 LOC
-

SLIMCACHE: CUCKOO HASHING



IN SUMMARY

- we spent months to achieve things we “already have”...
 - ... with more predictable resource, better logging/metrics...
 - ... out of a framework we can own long-term...
 - ... easy to add stuff to- the fun part!
-

NOW WHAT?



FOR TWITTER

- drop-in replacement for all in-house backend
 - new protocol and/or new features*
 - “unified cache”: build & migrate (& profit)
-

FOR THE REST

- we are open-sourcing! public in 2-3 weeks @pelikan_cache
 - serious about OSS, develop-in-the-open kind of seriousness
 - what is on your wishlist for cache?
-

LESSONS LEARNED

- funding this type of project was hard
 - refactoring is a continual process- nothing is sacred
 - form *influences* function
 - forward-thinking, but be prepared for predictions to be wrong
 - consistency is key to style
-

